

Appendix

A	benchmark	10	987
A.1	Statistics	10	988
B	Data Quality Inspection Details	10	989
B.1	Quality Control Mechanisms	10	990
B.2	Human Performance Baseline	10	991
B.3	Annotator Recruitment and Compensation	10	992
C	Experiments	11	993
D	Metric Exploration and Phenomenon Analysis	11	994
D.1	Evaluation Metric	11	995
D.2	Design Rationale	11	996
D.3	Formal Definition	12	997
D.4	Interpretation	12	998
D.5	Why Not Standard F1?	12	999
D.6	Language-Specific Results for Task 3 and Task 4	13	1000
D.7	Generative Evaluation for Task 3 and Task 4	13	1001
E	Discussion	17	1002
E.1	Impact of Irrelevant Charts	17	1003
E.2	Visual Dependency and Language Priors	17	1004
E.3	Performance Across Different Languages	17	1005
E.4	Exploring MLLMs' Retrieval Capabilities	17	1006
E.5	Error Analysis	17	1007
F	Extended Benchmark.	17	1008
F.1	Task Description	17	1009
F.2	Pipeline	20	1010
F.3	Data Statistics	21	1011
G	QAs Prompts	21	1012
			1013
			1014
			1015
			1016
			1017
			1018
			1019
			1020
			1021
			1022
			1023
			1024
			1025
			1026
			1027
			1028
			1029
			1030
			1031
			1032
			1033
			1034
			1035
			1036
			1037
			1038
			1039
			1040
			1041
			1042
			1043
			1044

A benchmark

A.1 Statistics

As summarized in Table 1 in the main text, each language version of MultiChartQA-R contains 180 multi-chart sets, 695 chart-code pairs, and 2,160 QA pairs, with 540 instances for each of the four tasks. On average, each set contains 3.9 charts, reflecting the realistic need to aggregate evidence across multiple visualizations rather than reason over isolated charts. The same construction pipeline is used for the English, Chinese, and Spanish versions, so the benchmark remains structurally aligned across languages while supporting multilingual evaluation.

Beyond scale, MultiChartQA-R is designed to provide broad diversity in both visual form and application context. Figure 3 shows that the benchmark covers 14 chart categories, while Figure 4 shows coverage across 36 domains. This diversity is important for reducing over-specialization to a narrow set of chart styles or topics, and for better reflecting the heterogeneous analytical scenarios encountered in real-world multi-chart reasoning.

Taken together, these statistics show that MultiChartQA-R is not only balanced in task composition, but also broad in chart forms and domains, making it suitable for studying both controlled task difficulty and generalization across realistic multi-chart scenarios.

B Data Quality Inspection Details

To ensure the reliability of MultiChartQA-R, we applied task-specific human verification throughout the dataset construction process and complemented it with a post hoc quality audit and a human performance benchmark.

B.1 Quality Control Mechanisms

Tasks 1 & 2: Full Cross-Review. For *Trend Inference* and *Data Integration*, which have objectively verifiable answers, we conducted full human cross-review over the complete dataset. The annotation team consisted of four members, and each item was independently checked by another annotator in a cyclic peer-review process. The review focused on three aspects:

- **Question validity:** whether the question strictly followed the task definition and was unambiguous;
- **Reasoning correctness:** whether the solution process was logically sound and consistent with the chart evidence;
- **Answer accuracy:** whether the final answer was fully supported by the chart data.

Any identified issues were corrected before the data were finalized.

Tasks 3 & 4: Human Refinement of Model-Assisted Synthesis. For *Anomaly and Pattern Analysis* and *Strategy Recommendation*, we adopted a model-assisted synthesis pipeline to generate candidate question-answer items. However, these outputs were not included directly in the benchmark. Instead, all synthesized samples were comprehensively reviewed and refined by human annotators before inclusion. The human revision process focused on four aspects:

- **Question-type alignment:** whether each question matched the intended task type and required the targeted form of multi-chart reasoning;
- **Validity of correct options:** whether the correct option and its explanation were well grounded in the chart evidence and logically coherent;
- **Effectiveness of distractors:** whether incorrect options were plausible yet ultimately invalid, including both easy negatives and hard negatives;
- **Explanation consistency:** whether the explanations for both correct and incorrect options were faithful to the chart evidence and consistent with the intended reasoning path.

This procedure ensured that the final Task 3 and Task 4 data were human-refined benchmark instances rather than raw model outputs.

Post Hoc Quality Audit for Tasks 3 & 4. After the full human refinement stage, we further conducted an expert quality audit on a randomly sampled 30% subset of Tasks 3 and 4. Evaluators scored each sample on a 10-point scale using the criteria above. The audit results indicate strong data quality: Task 3 achieved an average score of **9.1/10** with an inter-rater agreement of **85%**, while Task 4 achieved an average score of **9.3/10** with an inter-rater agreement of **87%**.

B.2 Human Performance Baseline

To establish a human-level reference and verify task difficulty, the expert group solved the benchmark without access to the ground truth. Human performance reached **97.83** and **94.83** on Task 1 and Task 2, **90.60** and **91.60** in the multi-select setting of Task 3 and Task 4, and **85.80** and **87.50** in the corresponding generative setting. Although human performance consistently exceeded that of current models, the more complex tasks still led to a non-negligible error rate, highlighting the intrinsic difficulty and research value of MultiChartQA-R.

B.3 Annotator Recruitment and Compensation

We recruited four annotators with undergraduate degrees or above in STEM-related fields to ensure sufficient expertise in chart interpretation and logical reasoning.

Compensation. For annotation, review, and human evaluation, annotators were paid at a rate of **\$15 per hour**, adjusted according to task difficulty. This rate is substantially higher than the local minimum wage (approximately **\$5 per hour**), ensuring fair compensation.

Consent. All annotators were informed of the purpose of the study and the intended use of the collected annotations. They provided consent for their contributions to be used for research purposes and released under the project license.




Figure 3: Distribution of Chart Types




Figure 4: Distribution of Domains

C Experiments

D Metric Exploration and Phenomenon Analysis

To validate the effectiveness of the proposed **Strict Risk-Aware MF_β Score**, we conducted a comparative analysis against standard metrics (Original F1) and recent benchmarks (COM2 from ACL 2025) on the English multi-select results of Task 3 and Task 4. The detailed results are presented in Table 4, where the **Strict** column is kept consistent with the main leaderboard.

The "Safety Gap" Phenomenon. A crucial observation from our experiments is the divergence between standard precision-based metrics and our risk-aware metric, which we term the "Safety Gap."

- **High-Performance Models:** For robust models such as Claude-Sonnet-4 and Seed1.5-VL, the gap between 'Orig F1' and 'Strict' is relatively moderate. For example, in Task 4 English, Claude drops from 87.53 to 77.49, a Δ of about 10.0. This indicates that these models not only identify correct answers but also avoid selecting high-risk distractors.
- **Low-Performance Models:** Conversely, weaker models such as DeepSeek-VL-7B exhibit a much larger drop (e.g., in Task 4 English, dropping from 56.95 to 29.32, a Δ of about 27.6).

This phenomenon reveals that standard metrics often inflate the performance of weaker models that engage in "reckless guessing" (hitting some correct options while simultaneously selecting fatal errors). The **Strict** metric successfully penalizes this behavior, offering a more distinct differentiation capability than 'Orig F1' and providing a balanced assessment compared to the overly harsh 'COM2' metric.

Correlation with Human Judgment. As shown in the table, the ranking consistency provided by the 'Strict' metric aligns more closely with the capability tiers observed in manual inspections. While 'Orig F1' might suggest that a 7B model performs comparably to a 72B model in certain scenarios due to high recall, the 'Strict' score exposes the fragility of the smaller model's reasoning, ensuring that high scores truly reflect reliable, deployable decision-making capabilities.

D.1 Evaluation Metric

The original MF_β score (Orig F1) and Com^2 exhibit similar distributions across model evaluations, suggesting that MF_β captures multiple-choice capability reasonably well. However, Com^2 amplifies the impact of distractors, which limits its applicability across scenarios. By contrast, MF_β is more flexible because it can adapt to different settings through the choice of β .

Anomaly and Pattern Analysis and Strategy Recommendation involve open-ended multi-chart question-answering, where the answers are not unique. These tasks are designed in a multiple-choice format with Easy and Hard distractors. To evaluate model performance, we propose the **Strict Risk-Aware MF_β** metric with fixed severity weights.

D.2 Design Rationale

The key motivation is that not all false positives are equally problematic in our benchmark. Some distractors are included as *easy errors*: these options either ignore the relevant chart evidence or lead to clearly incorrect conclusions. Selecting such an option indicates a basic failure of chart grounding or decision reliability. By

Model	Task 3 (Anomaly Analysis) - EN			Task 4 (Strategy Rec) - EN		
	Orig F1	Strict	COM2	Orig F1	Strict	COM2
<i>Proprietary Models</i>						
Claude-Sonnet-4	84.92	72.33	65.04	87.53	77.49	69.67
Gemini-2.5-Pro	81.87	74.66	59.04	83.79	75.54	60.86
Seed1.5-VL	78.87	74.79	58.59	82.22	78.40	66.15
GPT-4o	71.76	52.06	38.59	76.88	56.16	38.86
<i>Open-Weight Models</i>						
InternVL3-78B	81.62	71.80	57.15	81.78	73.00	57.20
Qwen3-VL-32B	84.64	67.80	63.29	86.63	68.35	64.45
Qwen2.5-VL-72B	72.62	60.02	32.93	76.34	63.57	39.44
InternVL2-L3-76B	68.33	55.26	34.07	70.19	58.54	36.95
Qwen2.5-VL-7B	71.38	58.32	37.33	74.55	62.21	43.19
MiniCPM-V-2_6	60.67	40.48	21.27	60.13	41.23	21.90
DeepSeek-VL-7B	49.90	24.61	11.35	56.95	29.32	17.11

Table 4: Comparison of different evaluation metrics on English multi-select tasks. Orig F1 represents the standard F1 score, which treats all errors equally. COM2 is a stringent metric from recent literature. Strict reports the same English Strict Risk-Aware MF_β scores as the main leaderboard, showing stronger discriminative power by penalizing models that make fatal errors more severely.

contrast, *hard errors* are intentionally designed as plausible near-miss alternatives: they may rely on partially correct evidence or otherwise resemble valid reasoning, but ultimately remain incorrect. We therefore penalize easy errors more heavily than hard errors.

Importantly, the fixed weights were determined from the benchmark’s annotation design before model-level analysis, rather than tuned to maximize any particular model ranking. Our goal is not to favor a specific system, but to reflect the benchmark’s intended risk structure: options that more clearly violate chart evidence should incur larger penalties.

D.3 Formal Definition

Let A be the set of correct answers (Ground Truth) and B be the set of answers selected by the model. The incorrect options are categorized into two subsets: Easy/Obvious Errors (E) and Hard/-Confusing Errors (H). We define:

$$TP = |A \cap B|, \quad e = |B \cap E|, \quad h = |B \cap H| \quad (5)$$

The **Net True Positive** with fixed severity weights is:

$$TP_{net}^* = \max(0, TP - w_e \cdot e - w_h \cdot h) \quad (6)$$

The fixed severity weights are:

$$w_e = 1.0, \quad w_h = 0.5 \quad (7)$$

Here, “easy” and “hard” refer to distractor discriminability rather than evaluation importance. The larger penalty on easy errors reflects the fact that these options correspond to clearer grounding failures or explicitly incorrect conclusions, whereas hard errors are designed to be more plausible alternatives.

Based on TP_{net}^* , we define:

$$P_{strict}^* = \begin{cases} 0, & |B| = 0 \\ \frac{TP_{net}^*}{|B|}, & |B| > 0 \end{cases} \quad R = \begin{cases} 0, & |A| = 0 \\ \frac{|A \cap B|}{|A|}, & |A| > 0 \end{cases} \quad (8)$$

The final $MF_\beta^{strict*}$ score is:

$$MF_\beta^{strict*} = \begin{cases} 0, & \beta^2 P_{strict}^* + R = 0 \\ (1 + \beta^2) \cdot \frac{P_{strict}^* R}{\beta^2 P_{strict}^* + R}, & \text{otherwise} \end{cases} \quad (9)$$

D.4 Interpretation

- When $\beta = 1$, the metric balances precision and recall equally.
- When $\beta > 1$, greater emphasis is placed on avoiding incorrect selections (higher precision).
- When $\beta < 1$, greater emphasis is placed on covering correct options (higher recall).

The *cancellation logic* in TP_{net}^* means that severe errors actively deduct from the model’s score rather than merely lowering precision. This mirrors high-stakes decision-making, where a single fatal flaw can render a strategy unusable.

D.5 Why Not Standard F_1 ?

Standard F_1 treats all false positives equally, regardless of whether the selected option is an obvious grounding failure or a plausible near-miss. This is unsuitable for our setting, where distractors are deliberately constructed with different semantic roles and different degrees of risk. The proposed metric preserves the familiar precision-recall structure of F_β while incorporating task-specific error severity through a simple and fixed penalty scheme. As shown in table 4, this makes the metric more discriminative than standard F_1 and more aligned with the benchmark’s decision-oriented evaluation objective.

We compare $MF_\beta^{strict*}$ with Com^2 [37] in section 3.2, where we highlight $MF_\beta^{strict*}$ ’s role in model selection for specific scenarios and its feasibility for preference model training.

Analysis of MF_β Curves

As shown in Figures 5 and 6, we plotted the MF_β curves of all models under varying values of β . Overall, the curves exhibit a monotonically increasing trend, indicating that within the current task setting, “selecting all correct options” is significantly more difficult than “avoiding incorrect options.” In other words, the models generally perform better at eliminating incorrect options than at fully covering the correct ones. We hypothesize that this phenomenon is related to the characteristics of the hard options: such options are more prone to being selected by the models, yet their penalty weight in the scoring scheme is relatively low, which to some extent “inflates” the overall scores. This effect is particularly pronounced for smaller-parameter models whose performance approaches random selection.

Intersection points. The intersections between the curves carry critical implications: they reveal the trade-offs between different models in terms of “recalling correct options” versus “avoiding incorrect ones,” thereby providing guidance for model selection across application scenarios. For example, in recall-oriented tasks, models that perform better before the intersection point should be prioritized, whereas in precision-oriented tasks, models that excel after the intersection point are preferable.

Curve variability. Taking InternVL2-26B as an example, its curve ranks relatively low when β is small, but improves markedly as β increases. This pronounced change highlights the model’s substantial variability across different evaluation emphases, reflecting a lack of balance—that is, an insufficiently stable ability to reconcile recall and precision.

D.6 Language-Specific Results for Task 3 and Task 4

Table 5 presents the detailed Task 3 (Anomaly/Pattern Analysis) and Task 4 (Strategy Recommendation) results across Chinese (cn), English (en), and Spanish (es) for the multi-select setting, together with generative scores kept consistent with the main leaderboard.

Key observations from this breakdown are as follows:

- **Cross-lingual variation in multi-select:** Multi-select performance on Task 3 varies across languages for several

models, indicating that anomaly and pattern analysis remains sensitive to language-specific reasoning conditions.

- **Task 4 strategic reasoning:** Strategy recommendation remains difficult in the multi-select setting, while the generative scores provide a complementary view that is directly aligned with the main leaderboard.
- **Open-weight competitiveness:** Qwen3-VL-32B and InternVL3-78B remain among the strongest open-weight models on Task 3 and Task 4, consistent with the main experiments.

D.7 Generative Evaluation for Task 3 and Task 4

To complement the multiple-choice setting, we additionally evaluate Task 3 (*Anomaly and Pattern Analysis*) and Task 4 (*Strategy Recommendation*) in a generative setting. In this setting, models are given the chart images and the question only, without access to the candidate options, and are required to produce a free-text analysis. Our goal is to evaluate both whether the generated analysis implies the correct decision and whether the accompanying reasoning is faithful, complete, and grounded in the chart evidence.

D.7.1 Evaluation Pipeline. The generative evaluation consists of four stages:



- **Stage 1: Generative inference.** The model receives chart images and the question, and produces a free-text analysis.
- **Stage 2: Answer extraction.** A judge LLM reads the generated analysis together with the original multiple-choice question and extracts the option(s) implied by the response.
- **Stage 3: Risk-aware option-level scoring.** The extracted option set is evaluated against the original annotations using the same Strict Risk-Aware MF_β protocol as in the multiple-choice setting.
- **Stage 4: Explanation-quality evaluation.** We further assess the generated explanation with an LLM judge and report reference-based similarity metrics as supplementary indicators.

This design keeps the primary generative score directly comparable to the original multiple-choice benchmark, while also providing a richer analysis of explanation quality.

D.7.2 Answer Extraction and Risk-Aware Scoring. Given a free-text response, we use a judge LLM to determine which option(s) the response implicitly endorses. The judge is provided with the original question and the generated analysis, and is asked to output the corresponding option letters. If the response does not clearly support any option, the judge outputs NONE.

Listing 1: Answer extraction prompt template

```
1 Given the following question and a candidate's
   textual analysis,
2 determine which option(s) the candidate is
   suggesting as the answer.
3
4 Question: {question}
5
6 Candidate's Analysis:
```

Figure 5: Model Performance on Task 3 under Different β Figure 6: Model Performance on Task 4 under Different β

Model	Task 3: Anomaly/Pattern Analysis				Task 4: Strategy Recommendation			
	Multi-select			Generative	Multi-select			Generative
	cn	en	es		cn	en	es	
Human	90.60	90.60	90.60	85.80	91.60	91.60	91.60	87.50
<i>Proprietary Models</i>								
Claude-Sonnet-4	75.78	70.00	69.33	77.6	77.49	77.49	77.49	73.1
Gemini-2.5-Pro	79.33	75.06	75.11	75.5	75.54	75.54	75.54	79.7
Seed1.5-VL	71.78	72.44	67.33	62.8	78.40	78.40	78.40	74.4
GPT-4o	62.64	64.21	59.87	58.2	56.16	56.16	56.16	73.2
<i>Open-weight Models</i>								
InternVL3-78B	68.46	73.21	64.66	64.2	73.00	73.00	73.00	61.8
Qwen3-VL-32B	73.94	73.11	72.67	76.3	68.35	68.35	68.35	77.8
Qwen3-VL-8B	64.44	62.36	60.44	64.2	70.75	70.75	70.75	73.7
InternVL2-L3-76B	48.88	59.91	59.19	62.9	53.51	58.54	51.14	69.7
Qwen2.5-VL-72B	56.95	56.25	53.13	70.7	63.57	63.57	63.57	76.3
LLaVA-OV-72B	57.59	61.33	53.33	54.8	52.85	52.85	52.85	58.4
InternVL2-26B	58.65	54.46	50.11	47.0	51.22	51.22	51.22	47.0
Qwen2.5-VL-7B	54.91	54.44	52.22	61.7	62.21	62.21	62.21	65.6
MiniCPM-V-2_6	55.36	49.88	44.22	54.1	41.23	41.23	41.23	59.6
DeepSeek-VL-7B	47.11	49.32	40.44	50.5	29.32	29.32	29.32	53.0
LLaVA-OV-7B	34.00	48.55	46.00	49.9	38.77	38.77	38.77	58.5

Table 5: Task 3 (Anomaly/Pattern Analysis) and Task 4 (Strategy Recommendation) results breakdown by language (cn, en, es). Multi-select columns report Strict Risk-Aware MF_β ($\beta = 1$, $w_e = 1.0$, $w_h = 0.5$), and Generative columns report the same free-form generation scores as the main leaderboard. Best scores in each column are in bold; “-” indicates unavailable results.

```

7 {generated_text}
8
9 Extract the option letter(s) (e.g., A, B, C, D
10 , E, F) that the
11 analysis implies as the correct answer.
12 If the analysis does not clearly indicate any
13 option, respond with:
14 NONE

```

After answer extraction, we evaluate the predicted option set against the original task annotations, including the correct labels as well as the *easy-error* and *hard-error* option groups. We then compute the same Strict Risk-Aware MF_β score as in the multiple-choice setting. In this way, the primary generative score remains aligned with the benchmark’s original risk-aware evaluation principle, rather than relying solely on free-form semantic similarity.

D.7.3 Explanation-Quality Evaluation. In addition to answer correctness, we evaluate the quality of the generated explanation with an LLM judge along four dimensions:

- **Evidence Relevance (0–1):** whether the cited chart evidence, values, and trends are directly relevant to the question and final conclusion;

- **Reasoning Completeness (0–1):** whether the reasoning chain covers the key intermediate steps without major omissions;
- **Hallucination Risk (0–1):** whether the explanation avoids introducing unsupported claims that are not grounded in the chart evidence; higher is better;
- **Consistency (0–1):** whether the intermediate statements, reasoning process, and final conclusion are mutually consistent.

The judge is prompted to output structured scores in JSON format:

Listing 2: Explanation-quality judge prompt template

```

1 You are an expert evaluator for multi-chart
2 reasoning explanations.
3 Evaluate the candidate's explanation based on
4 the following rubric.
5 Output strictly as JSON.
6 Question: {question}
7 Gold Correct Answer: {gold_answer}
8 Gold Explanation: {gold_explanation}
9

```

```

1741 10  Candidate's Analysis:
1742 11  {candidate_analysis}
1743 12
1744 13  Rubric:
1745 14  - evidence_relevance: Whether cited evidence
1746    is directly relevant
1747 15  to the question and final answer.
1748 16  - reasoning_completeness: Whether the
1749    reasoning chain covers key
1750 17  steps without major omissions.
1751 18  - hallucination_risk: Whether the explanation
1752    introduces unsupported
1753 19  facts not grounded in chart evidence.
1754 20  - consistency: Whether reasoning, intermediate
1755    statements, and final
1756 21  answer are mutually consistent.
1757 22
1758 23  Output format:
1759 24  {
1760 25    "evidence_relevance": float (0-1),
1761 26    "reasoning_completeness": float (0-1),
1762 27    "hallucination_risk": float (0-1),
1763 28    "consistency": float (0-1),
1764 29    "overall": float (0-1),
1765 30    "comments": str
1766 31  }

```

This stage is intended to capture explanation quality beyond discrete option correctness.

D.7.4 Reference-Based Similarity. We further report reference-based similarity metrics between the generated explanation and the gold explanation, including **BLEU-4**, **BERTScore**, and **Rouge-L**. These metrics serve as supplementary indicators of lexical and semantic overlap, but are not used as the primary leaderboard metric, since a valid explanation may differ substantially in wording from the reference while still being correct and well grounded.

D.7.5 Reported Metrics. In the main leaderboard, we report **Strict F1** as the primary score for generative evaluation, where the predicted answer set is obtained through answer extraction and then evaluated under the same risk-aware option-level protocol as in the multiple-choice setting. We additionally report the four LLM-judge dimensions and the reference-based similarity metrics in supplementary analyses to provide a more complete picture of generative reasoning quality.

E Discussion

E.1 Impact of Irrelevant Charts

In these experiments, because the number of charts varies and high-quality charts with many visual tokens can cause the total input length to exceed the model's context window, the model will begin discarding tokens from the left (i.e., the earliest inputs) until the remaining token count less than `session_len`. Therefore, we deliberately placed any image tokens for charts unrelated to the QA pair on the far left; even if they are dropped, the model's ability to answer correctly remains unaffected.

Method for handling excessively long image tokens. In these experiments, because the number of charts varies and high-quality charts with many visual tokens can cause the total input length to exceed the model's context window, the model will begin discarding tokens from the left (i.e., the earliest inputs) until the remaining token count less than `session_len`. Therefore, we deliberately placed any image tokens for charts unrelated to the QA pair on the far left; even if they are dropped, the model's ability to answer correctly remains unaffected.

E.2 Visual Dependency and Language Priors

To rigorously verify whether the models rely on visual perception or exploit textual shortcuts, we conducted a blind ablation study (Table 8) by removing the chart inputs. The results highlight three key findings:

Necessity of Visual Information. In the *Text Only* setting of Task 1 (Trend Inference), models achieve scores around 47%–50%, effectively performing random guessing on the binary task. However, when explicitly instructed to answer "unknown" if charts are missing (*Text+Hint*), accuracy drops precipitously (e.g., Seed1.5-VL drops to 3.56%). Since the ground truth labels are binary (Yes/No), this near-zero score confirms that models successfully ceased random guessing and followed instructions to reject the query. This sharp contrast validates that the high performance in the *Standard* setting is driven by genuine visual perception rather than textual speculation. Similarly, the collapse of performance in Task 2 (Data Integration) without images (e.g., InternVL3-78B: 5.60%) further proves that precise numerical reasoning in our benchmark is strictly vision-dependent.

LLM Priors vs. Visual Grounding. In Task 4 (Strategy Recommendation), models maintain relatively high baselines even in the *Text Only* setting (e.g., Seed1.5-VL: 76.04%). This suggests that advanced LLMs can leverage their internal knowledge bases to propose plausible business strategies based solely on the textual context of the question. However, introducing visual information (*Standard*) consistently yields performance gains (increasing to 78.46%). This demonstrates that while language priors provide a strong foundation for decision-making, visual alignment is essential for tailoring strategies to the specific data trends presented in the charts, thereby achieving the optimal solution.

E.3 Performance Across Different Languages

Robustness in Mixed-Language Scenarios. The results on mixed-language charts (Table 13) further support the "language-agnostic" nature of advanced visual perception. Top-tier models such as




Figure 7: Compare the performance of MLLMs across four retrieval tasks. The x-axis represents the number of charts involved in the question, for example, 2c denotes two charts.

Seed1.5-VL and InternVL3-78B maintain performance close to their monolingual baselines even when English, Chinese, and Spanish charts are interleaved within a single query. For example, Seed1.5-VL achieves 70.38% on Trend Inference with mixed charts, close to its English-only baseline of 72.44%. This suggests that SOTA models handle multi-chart QA not by relying on one unified linguistic context, but by performing semantic extraction on each chart regardless of source language, showing strong generalization in heterogeneous real-world settings.

E.4 Exploring MLLMs' Retrieval Capabilities

E.5 Error Analysis

The detailed response statistics of the four proprietary models on Task 3 and Task 4 are presented in Tables 14 & 15.

F Extended Benchmark.

F.1 Task Description

parallel type.

- EXAMPLE

question: one sub-question for chart 1; one for chart 2; one for chart 3; and one for chart 4

answer: answer1. answer2. answer3. answer4.

Model	Trend Inference		Data Integration		Anomaly/Pattern Attr		Strategy Rec	
	involved	all	involved	all	involved	all	involved	all
Claude-Sonnet-4	70.00	69.11	60.99	61.38	73.99	73.52	78.21	78.31
Seed1.5-VL	72.44	72.83	67.66	66.74	68.85	68.26	73.84	71.50
GPT-4o	64.21	66.14	64.83	63.10	49.05	47.18	48.96	49.08
InternVL3-78B	73.21	60.93	67.50	64.68	66.28	65.68	65.68	63.05
InternVL2-Llama3-76B	59.91	52.46	51.59	44.36	46.38	45.22	47.24	46.86
Qwen2.5-VL-72B-Instruct	56.25	58.93	25.40	22.94	48.23	50.33	52.18	48.93
LLaVA-OV-72B	61.33	59.11	43.02	38.07	40.76	40.29	43.21	42.73
InternVL2-26B	54.46	46.33	31.03	22.90	47.60	42.82	48.96	48.85
Qwen2.5-VL-7B-Instruct	44.00	53.56	21.04	23.98	51.93	45.11	53.30	49.28
MiniCPM-V-2_6	49.88	52.89	23.29	27.48	33.81	35.85	30.93	32.78
DeepSeek-VL-7B	49.32	24.22	7.95	3.48	21.23	16.19	24.09	23.87
LLaVA-OV-7B	48.55	45.19	21.84	13.55	19.15	19.81	22.88	19.74

Table 6: Compare the question-answering performance when inputting all charts versus only the relevant charts.

Model	Trend Inference		Data Integration		Anomaly/Pattern Attr		Strategy Rec	
	involved	all	involved	all	involved	all	involved	all
Claude-Sonnet-4	70.00	69.11	60.99	61.38	65.04	62.56	69.67	69.15
Seed1.5-VL	72.44	72.83	67.66	66.74	58.72	58.63	66.30	63.81
GPT-4o	64.21	66.14	64.83	63.10	39.29	35.55	39.13	38.54
InternVL3-78B	73.21	60.93	67.50	64.68	57.37	54.25	57.20	54.20
InternVL2-Llama3-76B	59.91	52.46	51.59	44.36	34.24	33.99	36.95	36.18
Qwen2.5-VL-72B-Instruct	56.25	58.93	25.40	22.94	32.93	33.19	39.44	37.26
LLaVA-OV-72B	61.33	59.11	43.02	38.07	27.26	27.48	31.67	30.91
InternVL2-26B	54.46	46.33	31.03	22.90	36.32	30.81	36.53	37.27
Qwen2.5-VL-7B-Instruct	44.00	53.56	21.04	23.98	37.33	32.81	43.19	39.52
MiniCPM-V-2_6	49.88	52.89	23.29	27.48	21.27	24.63	21.90	22.86
DeepSeek-VL-7B	49.32	24.22	7.95	3.48	11.35	7.59	17.19	16.30
LLaVA-OV-7B	48.55	45.19	21.84	13.55	9.20	8.97	12.03	10.13

Table 7: Compare the question-answering performance when inputting all charts versus only the relevant charts under the com2 metric.

Model	Trend Inference			Data Integration		Anomaly/Pattern Attr		Strategy Rec	
	Text+Hint	Text Only	Standard	Text Only	Standard	Text Only	Standard	Text Only	Standard
Seed1.5-VL	3.56	47.33	71.78	6.31	72.81	63.76	75.35	76.04	78.46
InternVL3-78B	14.51	47.44	68.46	5.60	70.72	51.46	66.38	53.94	68.24
MiniCPM-V-2_6	21.56	46.00	55.36	4.66	26.15	26.90	31.44	27.33	35.02

Table 8: Ablation study on visual modality dependency. Standard: The model receives both chart images and text prompts. Text Only: The model receives only text prompts without chart images. Text+Hint: In Task 1, the model receives only text but is explicitly instructed to output "unknown" if visual information is missing. Note that in the 'Text+Hint' setting, an output of "unknown" is counted as incorrect (0 score) against the ground truth (Yes/No), strictly penalizing non-visual guessing.

Model	Trend Inference			Data Integration			Anomaly/Pattern Attr			Strategy Rec		
	en	zh	es	en	zh	es	en	zh	es	en	zh	es
Seed1.5-VL	72.44	71.56	67.56	67.66	68.76	67.04	68.85	72.78	68.80	73.84	77.52	71.41
InternVL3-78B	73.21	70.98	66.00	67.50	70.69	65.91	66.28	68.10	62.04	65.68	67.32	62.48
InternVL2-26B	54.46	46.09	45.07	31.03	28.57	21.72	47.60	38.85	37.54	48.96	42.59	39.82
MiniCPM-V-2_6	49.88	51.79	48.78	23.29	29.50	23.95	33.81	36.99	27.53	30.93	35.12	28.61
Qwen2.5-VL-7B	54.44	55.23	48.67	21.04	23.81	23.29	51.93	48.29	50.82	53.30	47.11	60.56
LLaVA-OV-7B	48.55	33.48	46.33	21.84	18.26	13.14	19.15	22.14	30.73	22.88	32.11	27.77

Table 9: English Charts - Different Language QAs

Model	Trend Inference			Data Integration			Anomaly/Pattern Attr			Strategy Rec		
	en	zh	es	en	zh	es	en	zh	es	en	zh	es
Seed1.5-VL	72.44	68.22	69.78	67.66	65.69	65.70	68.85	71.86	70.75	73.84	74.36	71.66
InternVL3-78B	73.21	68.97	69.87	67.50	67.26	58.03	66.28	66.24	64.55	65.68	65.62	66.36
InternVL2-26B	54.46	53.56	48.44	31.03	18.20	19.46	47.60	46.80	42.95	48.96	53.25	47.83
MiniCPM-V-2_6	49.88	50.00	49.56	23.29	24.53	25.29	33.81	36.41	36.05	30.93	31.79	30.78
Qwen2.5-VL-7B	54.44	52.35	52.67	21.04	14.41	20.95	51.93	48.70	53.74	53.30	49.16	53.56
LLaVA-OV-7B	48.55	41.67	46.85	21.84	5.11	14.25	19.15	18.75	19.92	22.88	17.95	20.32

Table 10: Different Language Charts - English QAs

Model	Trend Inference			Data Integration			Anomaly/Pattern Attr			Strategy Rec		
	en	zh	es	en	zh	es	en	zh	es	en	zh	es
Seed1.5-VL	72.44	71.56	67.56	67.66	68.76	67.04	58.72	63.57	59.39	66.30	69.63	62.52
InternVL3-78B	73.21	70.98	66.00	67.50	70.69	65.91	57.15	55.21	51.69	57.20	58.60	52.78
InternVL2-26B	54.46	46.09	45.07	31.03	28.57	21.72	36.32	26.54	26.76	36.53	31.22	29.44
MiniCPM-V-2_6	49.88	51.79	48.78	23.29	29.50	23.95	21.27	25.93	17.44	21.90	24.93	18.01
Qwen2.5-VL-7B	54.44	55.23	48.67	21.04	23.81	23.29	37.33	35.05	39.07	43.19	36.07	48.67
LLaVA-OV-7B	48.55	33.48	46.33	21.84	18.26	13.14	9.20	11.19	18.52	12.03	20.04	17.04

Table 11: English Charts - Different Language QAs Under the Com2 Metric

Model	Trend Inference			Data Integration			Anomaly/Pattern Attr			Strategy Rec		
	en	zh	es	en	zh	es	en	zh	es	en	zh	es
Seed1.5-VL	72.44	68.22	69.78	67.66	65.69	65.70	58.72	63.06	61.43	66.30	66.00	63.52
InternVL3-78B	73.21	68.97	69.87	67.50	67.26	58.03	57.15	56.34	53.75	57.20	55.47	57.23
InternVL2-26B	54.46	53.56	48.44	31.03	18.20	19.46	36.32	35.51	30.54	36.53	41.07	35.94
MiniCPM-V-2_6	49.88	50.00	49.56	23.29	24.53	25.29	21.27	23.19	23.04	21.90	21.86	21.41
Qwen2.5-VL-7B	54.44	52.35	52.67	21.04	14.41	20.95	37.33	35.26	41.41	43.19	39.20	43.52
LLaVA-OV-7B	48.55	41.67	46.85	21.84	5.11	14.25	9.20	8.43	8.75	12.03	8.48	11.15

Table 12: Different Language Charts - English QAs Under the Com2 Metric

- PCPC

A question contains multiple parallel sub-questions, each retrieving one piece of information from a single chart.

Model	Trend Inference	Data Integration	Anomaly/Pattern Attr	Strategy Rec
Seed1.5-VL	70.38	65.11	69.54	74.54
InternVL3-78B	69.93	64.96	68.23	66.18
Qwen2.5-VL-7B-Instruct	54.12	19.38	49.94	51.17

Table 13: The test results of English Q&A on mixed-language charts.

Option	claude	seed	gpt	gemini
A	24.5 / 6.3	36.8 / 0.8	64.2 / 1.3	26.9 / 1.7
B	7.9 / 9.2	10.2 / 8.1	3.4 / 48.0	2.3 / 19.0
C	9.6 / 10.3	6.9 / 12.6	4.8 / 53.4	2.7 / 24.4
D	17.3 / 6.6	36.1 / 0.8	76.0 / 0.8	36.1 / 1.2
E	9.4 / 11.6	9.9 / 10.5	4.7 / 51.6	1.6 / 23.3
F	8.9 / 5.6	29.7 / 1.6	63.4 / 1.6	34.7 / 1.6
G	12.4 / 4.3	33.8 / 1.6	57.9 / 1.0	26.9 / 0.7

Table 14: Omission and multiple-selection rates per option for each model for task 3 (%). In this context, “claude” refers to Claude-Sonnet-4, “seed” refers to Seed1.5-VL, “gpt” refers to GPT-4o, and “gemini” refers to Gemini-2.5-Pro.

Option	claude	seed	gpt	gemini
A	11.6 / 4.4	40.7 / 1.2	62.8 / 0.4	31.6 / 0.5
B	4.3 / 9.1	2.9 / 7.9	1.9 / 49.2	3.0 / 21.1
C	2.5 / 6.8	2.5 / 10.4	1.0 / 47.4	1.2 / 30.5
D	8.9 / 2.7	32.8 / 0.0	62.0 / 1.2	25.2 / 0.9
E	4.2 / 11.6	7.8 / 12.8	2.1 / 58.1	0.0 / 32.7
F	9.1 / 4.4	25.6 / 0.7	58.0 / 2.9	26.3 / 3.8
G	7.3 / 4.0	35.8 / 0.0	58.9 / 2.3	34.2 / 0.0

Table 15: Omission and multiple-selection rates per option for each model for task 4 (%). In this context, “claude” refers to Claude-Sonnet-4, “seed” refers to Seed1.5-VL, “gpt” refers to GPT-4o, and “gemini” refers to Gemini-2.5-Pro.

2c, 3c, and 4c refer to retrieving one piece of information from charts 2, 3, and 4, respectively, and providing separate answers.

- PCMC
A question contains multiple parallel sub-questions, each retrieving multiple pieces of information from a single chart.
2c, 3c, and 4c refer to retrieving at least one piece of information from charts 2, 3, and 4, respectively, and providing separate answers.

union type.

- EXAMPLE
question: question about chart_1&2&3&4
answer: answer.
- PCPC
The question involves multiple charts, retrieving one piece of information from each chart, combining the information to perform a simple calculation, and outputting the final result.

- PCMC

The question involves multiple charts, retrieving at least one piece of information from each chart, combining the information to perform a simple calculation, and outputting the final result.

F.2 Pipeline

1. Manually select articles from Pew that contain at least four charts, and scrape both the articles and charts.
2. Extract chart information and perform manual sampling-based screening to ensure quality.
3. Automatically generate topic summaries based on the content of the article.
4. Generate question-and-answer pairs by utilizing the extracted chart information and setting different prompts based on the topics.
5. Use scripts to check whether the number of charts involved in all question-and-answer pairs and the amount of content related to each chart meet the expected criteria. If they do not meet the criteria, regenerate the pairs and manually

Category	Number of Articles
Economy & Work	11
Politics & Policy	11
Internet & Technology	11
Family & Relationships	11
Age & Generations	11
Immigration & Migration	11
Science	11
News Habits & Media	12
Other Topics	12

Table 16: Category-wise Number of Articles

adjust them until all question-and-answer pairs satisfy the conditions.

6. Use scripts to check the length of all answers, manually screen answers that exceed the threshold, remove redundant parts, and retain concise answers.

7. Utilize the rationale of generated question-and-answer pairs to generate code-based computation results, manually compare all inconsistent answers, correct the answers and rationales in the labels, and ensure the accuracy of the computation results.

F.3 Data Statistics

We collected 101 articles from Pew Research Center, each containing at least four charts centered on the same topic, spanning nine topics in total (Refer to Table 16), yielding 1,212 QA pairs. These QA pairs span four distinct question types, comprehensively probing models’ capabilities in cross-chart information retrieval.

G QAs Prompts

For Task 3 and Task 4, we use neutral multiple-choice prompt templates that do not contain one-shot answer demonstrations or fixed label patterns, in order to reduce prompt-induced option bias.

Task 1: Trend Inference

```

1 output_prompt = '''
2     Output according to the following json format:
3     {
4         "rationale": "Reasoning process",
5         "answer": "Final answer"
6     }
7     Please strictly output according to the json format.
8 '''
9 prompt = image_tokens + question + output_prompt

```

Task 2: Data Integration

```

1 output_prompt = '''
2     Please answer according to the following steps:
3         1. Extract relevant information...
4         2. Explain the association logic...
5         3. Calculation process...
6         4. Conclusion summary...
7         5. Output requirements...
8 '''
9 prompt = image_tokens + question + output_prompt

```

Task 3: Anomaly/Pattern Attribution

```

1 text1 = "The options for the question are as follows: "
2 hints = "(Multiple choice question) Please analyze..."
3 output_prompt = '''
4     Answer the question according to the following JSON format:
5     {
6         "rationale": "Reasoning process grounded in the chart evidence",
7         "answer": ["option_letter_1", "option_letter_2"]
8     }
9     The "answer" field should contain only the option letters you judge to be correct.
10    Do not imitate any fixed option pattern; determine the answer solely from the chart
11    content.
12 '''
13 prompt = image_tokens + hints + question + text1 + answer_choices + output_prompt

```

Task 4: Strategy Recommendation

```
1 text1 = "The options for the question are as follows: "  
2 hints = "(Multiple choice question) Please analyze the content of the chart and select the  
    correct options based on the chart information. Note: Correct options should be supported  
    by chart data; otherwise, they will be regarded as incorrect."  
3 output_prompt = '''\nAnswer the question according to the following JSON format:  
4     {  
5         "rationale": "Reasoning process grounded in the chart evidence",  
6         "answer": ["option_letter_1", "option_letter_2"]  
7     }  
8     Select the correct option letter(s) according to the actual chart content.  
9     Do not follow any fixed answer pattern; use only chart-supported evidence.  
10 '''  
11 prompt = image_tokens + hints + question + text1 + answer_choices + output_prompt
```

The prompt for generating python calculation code based on the reasoning process in task 2 is as follows:

```

1 gen_code_prompt = '''
2     The above is the reasoning process. Since large models are not good at calculations,
      ignore the calculated results in the reasoning process above, and generate executable
      Python code for the reasoning process.
3
4     Please note:
5     1. Strictly generate Python code based on the "rationale" process above.
6     2. Ensure the code correctly reflects the reasoning of the "rationale" by using variables
      to replace the intermediate calculation results in the rationale, as the code
      execution is more accurate and avoids cumulative errors caused by using intermediate
      calculation results.
7     3. The final output should strictly follow the required format, only printing the last
      output from the final `print` statement, without any descriptive print statements.
      Below are some examples of the final print outputs for reference:
8     a. Output format: "The answer should be presented as an integer."
      Incorrect final output 1: print(f"The final answer is: {answer}.")
      Incorrect final output 2: print(f"**{answer}**")
      Correct final output: print(f"{answer}")
9     b. Output in percentage, rounded to 4 significant digits.
      Incorrect final output 1: print(f"In mobile search, the proportion of images is
      **{answer}%**.")
      Correct final output: print(f"{answer}%")
10    c. Output in dollars, rounded to 3 decimal places. Answer: xx dollars.
      Incorrect final output 1: print(f"In Q1 2023, the in-app purchase revenue per
      download for mobile games on the App Store in Japan is {answer} USD.")
      Correct final output: print(f"{answer} USD")
11    d. Output "Yes" or "No".
      Incorrect final output 1: print("True")
      Correct final output: print("Yes")
12    e. Answer: xx times.
      Incorrect final output 1: print(f"{answer}")
      Correct final output: print(f"{answer} times")
13
14    4. Be sure to distinguish between "significant digits" and "decimal places." "Significant
      digits" refer to all digits from the first non-zero digit to the last digit. "Decimal
      places" refer to the number of digits after the decimal point.
15
16    5. If floating-point calculations are involved, please use the `decimal` library to avoid
      precision loss from floating-point operations.
17
18    '''
19
20
21
22
23
24
25

```